

# Modeling and Control of Bittide Synchronization



Sanjay Lall, Google Research and Stanford University


Călin Cașcaval, Martin Izzard, and Tammo Spalink, Google Research

# Overview

- *Google research project*
  - originated at Princeton: Spalink 2006.
  - broad scope: applications, scheduling, simulation, hardware, theory
- *paper*
  - problem formulation
  - model well-posedness
  - simulation algorithm
- *outline*: the mechanism, logical synchrony, controlling frequency


# Overview

- at each node  $i$ 
  - clock
  - each incoming link has a queue called the *elastic buffer*
- at each node, with each clock tick
  - a frame is removed from all elastic buffers
  - a frame is sent on all outgoing links




# Mechanism

- if oscillator at node  $j$  is faster than that at  $i$ 
  - $j$ 's elastic buffer will drain
  - $i$ 's elastic buffer will fill
- nodes observe elastic buffers, adjust frequency



# Logical Synchrony



- marked frames from nodes 1 and 3 always arrive simultaneously at node 2
- an example of *logical synchrony*
- does not require clocks to be synchronized

# Abstract frame model

# Modeling frames and phase

$$\beta_{i \rightarrow j}(t) = \lfloor \theta_i(t - l_{i \rightarrow j}) \rfloor - \lfloor \theta_j(t) \rfloor + \lambda_{i \rightarrow j}$$

- history of clock phases  $\theta$  determines location of every frame
- $\lambda_{i \rightarrow j}$  is a constant, determined by clock offsets at boot
- buffer occupancy is (roughly) phase difference between clocks at each end of the link

# Control loop

- *dynamics*

$$\frac{d\theta_i}{dt} = \omega_i$$
$$\beta_{i \rightarrow j}(t) = \lfloor \theta_i(t - l_{i \rightarrow j}) \rfloor - \lfloor \theta_j(t) \rfloor + \lambda_{i \rightarrow j}$$

- *measurements*

- at each node  $i$ , measure buffer occupancies  $\beta_{j \rightarrow i}$  for all neighbours  $j$  (relative to a mid-point value  $\beta^{\text{off}}$ )

- let  $r_i$  be the sum  $r_i = \sum_{j|j \sim i} (\beta_{j \rightarrow i} - \beta^{\text{off}})$

- *control*

- apply *frequency correction*:  $\omega_i = c_i + \omega_i^u$
- $\omega_i^u$  is the *uncorrected frequency* of the oscillator, not known
- e.g., proportional control is  $c_i = k_P r_i$



# The abstract frame model

for  $k \geq 0$  and  $t \in \mathbb{R}$

$$\dot{\theta}_i(t) = \omega_i(t)$$

$$\beta_{j \rightarrow i}(t) = \lfloor \theta_j(t - l) \rfloor - \lfloor \theta_i(t) \rfloor + \lambda_{j \rightarrow i}$$

$$\omega_i(t) = \omega_i^k \quad \text{for } t \in [s_i^k, s_i^{k+1})$$


$$t_i^k = \theta_i^{-1}(\theta_i^0 + kp)$$

$$s_i^k = \theta_i^{-1}(\theta_i^0 + kp + d)$$

$$\omega_i^k = c_i^k + \omega_i^u$$

$$c_i^k = \chi_i^k(m_i^0, m_i^1, \dots, m_i^k)$$

$$m_i^k = \{\beta_{j \rightarrow i}(t_i^k) \mid j \in \mathcal{V}, j \sim i\}$$



# The abstract frame model

for  $k \geq 0$  and  $t \in \mathbb{R}$

$$\dot{\theta}_i(t) = \omega_i(t)$$

$$\beta_{j \rightarrow i}(t) = \lfloor \theta_j(t - l) \rfloor - \lfloor \theta_i(t) \rfloor + \lambda_{j \rightarrow i}$$

$$\omega_i(t) = \omega_i^k \quad \text{for } t \in [s_i^k, s_i^{k+1})$$

$$t_i^k = \theta_i^{-1}(\theta_i^0 + kp)$$

$$s_i^k = \theta_i^{-1}(\theta_i^0 + kp + d)$$

$$\omega_i^k = c_i^k + \omega_i^u$$

$$c_i^k = \chi_i^k(m_i^0, m_i^1, \dots, m_i^k)$$

$$m_i^k = \{\beta_{j \rightarrow i}(t_i^k) \mid j \in \mathcal{V}, j \sim i\}$$

- features:

- sample rate at node  $i$  depends on  $\omega_i$
- no absolute time (integration/differentiation scaled by local clock rate)
- hybrid: quantization

- existence of solutions? yes, see paper

- exact algorithm for simulation


- software: [bittide.googleusercontent.com](http://bittide.googleusercontent.com)

# Control


# Control objectives

- frequencies cannot remain different for too long, otherwise buffers will over/underflow
- bittide performance requirement: maintain buffer occupancy within limits
- ideally buffer occupancies should be small, and frequencies large
- controller must be decentralized
- no *in-band* signalling
- failure handling, addition and removal of nodes, boot, *etc.*,

# Proportional control



- $k_P = 0.01$
- $\text{latency} = 1.0$
- $\text{poll\_period} = 10$
- $\text{uncorrected\_frequency} = (1.1, 1.4, 2.0)$
- $\text{control\_delay} = 2$
- varying step width is a consequence of periodic sampling w.r.t. the local clock



# Simpler models

- fluid model

$$\dot{\bar{\theta}} = \omega$$

$$\dot{\bar{\beta}} = -B^T \bar{\theta}$$


$$r = B \bar{\beta}$$

$$\omega = c + \omega^u$$

- with proportional control

$$\dot{\bar{\theta}} = -k_P B B^T \bar{\theta} + \omega^u$$

the well-known Laplacian dynamics



Thank you!