


On Buffer Centering for Bittide Synchronization



Sanjay Lall, Google Research and Stanford University


Călin Cașcaval, Martin Izzard, and Tammo Spalink, Google Research

Overview

- *Google research project*
 - originated at Princeton: Spalink 2006.
 - broad scope: applications, scheduling, simulation, hardware, theory
- *paper*
 - problem formulation
 - model well-posedness
 - simulation algorithm
- *outline:* the mechanism, logical synchrony, controlling frequency


Overview

- at each node i
 - clock
 - each incoming link has a queue called the *elastic buffer*
- at each node, with each clock tick
 - a frame is removed from all elastic buffers
 - a frame is sent on all outgoing links




Mechanism

- if oscillator at node j is faster than that at i
 - j 's elastic buffer will drain
 - i 's elastic buffer will fill
- nodes observe elastic buffers, adjust frequency



Logical Synchrony



- marked frames from nodes 1 and 3 always arrive simultaneously at node 2
- an example of *logical synchrony*
- does not require clocks to be synchronized

Abstract frame model

Modeling frames and phase

$$\beta_{i \rightarrow j}(t) = \lfloor \theta_i(t - l_{i \rightarrow j}) \rfloor - \lfloor \theta_j(t) \rfloor + \lambda_{i \rightarrow j}$$

- history of clock phases θ determines location of every frame
- $\lambda_{i \rightarrow j}$ is a constant, determined by clock offsets at boot
- buffer occupancy is (roughly) phase difference between clocks at each end of the link

Control loop

- *dynamics*

$$\frac{d\theta_i}{dt} = \omega_i$$
$$\beta_{i \rightarrow j}(t) = \lfloor \theta_i(t - l_{i \rightarrow j}) \rfloor - \lfloor \theta_j(t) \rfloor + \lambda_{i \rightarrow j}$$

- *measurements*

- at each node i , measure buffer occupancies $\beta_{j \rightarrow i}$ for all neighbours j (relative to a mid-point value β^{off})

- let r_i be the sum $r_i = \sum_{j|j \sim i} (\beta_{j \rightarrow i} - \beta^{\text{off}})$

- *control*


- apply *frequency correction*: $\omega_i = c_i + \omega_i^u$
- ω_i^u is the *uncorrected frequency* of the oscillator, not known
- e.g., proportional control is $c_i = k_P r_i$

Control


Control objectives

- frequencies cannot remain different for too long, otherwise buffers will over/underflow
- bittide performance requirement: maintain buffer occupancy within limits
- ideally buffer occupancies should be small, and frequencies large
- controller must be decentralized
- no *in-band* signalling
- failure handling, addition and removal of nodes, boot, etc.,

Proportional control






- $k_P = 0.01$
- $\text{latency} = 1.0$
- $\text{poll_period} = 10$
- $\text{uncorrected_frequency} = (1.1, 1.4, 2.0)$
- $\text{control_delay} = 2$
- varying step width is a consequence of periodic sampling w.r.t. the local clock



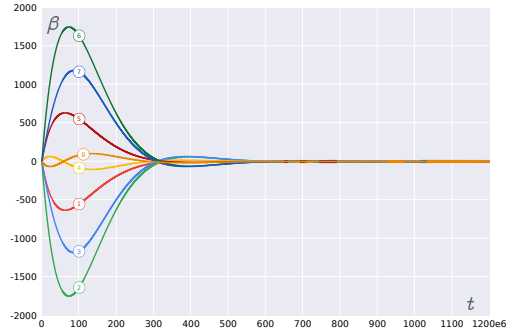


Example: Proportional control

- frequency correction is proportional to sum of relative buffer offsets
- equilibrium buffer offsets nonzero



Proportional-integral control

- ensures small steady state relative buffer occupancy



Reset control

- first, apply proportional control

$$c_i = k_P \sum_{j \rightarrow i} \beta_{i \rightarrow j}^{\text{rel}}$$

- this converges $c_i(t) \rightarrow c_i^{\text{ss}}$
- now change to proportional-plus-offset




$$c_i = c_i^{\text{ss}} + k_P \sum_{j \rightarrow i} \beta_{i \rightarrow j}^{\text{rel}}$$

- immediately after the switch, nodes are at the wrong frequency
- system is stable, so will return to equilibrium
- equilibrium happens when $c_i = c_i^{\text{ss}}$
- if the system returns to *the same* equilibrium, we will have

$$\sum_{j \rightarrow i} \beta_{i \rightarrow j}^{\text{rel}} \rightarrow 0$$

Reset

- controller shutdown at time $t \approx 4e9$
- at time $t \approx 6e9$, turn on the controller with new offset



Soft reset

- first, apply proportional control

$$c_i = k_P \sum_{j \rightarrow i} \beta_{i \rightarrow j}^{\text{rel}}$$




- this converges $c_i(t) \rightarrow c_i^{\text{ss}}$
- now *slowly* change to proportional-plus-offset

$$c_i = f(t)c_i^{\text{ss}} + k_P \sum_{j \rightarrow i} \beta_{i \rightarrow j}^{\text{rel}}$$

where $f(t)$ slowly changes from 0 to 1 over some interval

Example: soft reset

- controller shutdown at time $t \approx 4e9$
- soft reset over approximate interval $[6e9, 1e10]$



Summary

- soft reset
 - use P control initially
 - shutdown after convergence
 - turn on P control plus offset
- retains stability, keeps buffers at (or near) midpoint
- does not require integral control, spanning tree or global coordination

Thank you!